

CSC 202: Intro to Machine Intelligence

Natural Language Processing Project

In our unit involving Natural Language Processing, we covered how to analyze text in Python using the *TextBlob* and *Spacy* modules. In this project, you will use natural language processing for text-mining to analyze one or more texts. **Be creative!**

Analyze one or more sets of texts by completing the requirements below. Example project ideas include

- Analyzing a text, such as an e-book, Wikipedia entry, news article, political speech, or song lyrics
- Comparing texts (such as different Wikipedia pages, different e-books, different chapters from the same book, different songs, etc)

The following resources may be useful:

Song lyrics: <https://www.lyrics.com/>

E-books: [Project Gutenberg](#)

Movie scripts :

- [The Internet Movie Script Database](#)
- [Awesome Film](#)
- [Screenplays for You](#)
- [Movie Scripts and Screenplays](#)

You must turn in a Jupyter Notebook that carries out the appropriate analysis. Your notebook should begin with your name and a one paragraph summary describing the main goals of your project (e.g., to find the most common words in an e-book). Each code cell should be preceded by a markdown cell that describes what the code does (as in the class notes). Your Notebook should be correctly formatted with an appropriate title, section headers, and subsection headers. [20 points]

1. Complete **one** of the following to collect data to analyze [25 points]:

- Retrieve data from at least one Wikipedia page (see Lab 8 for an example)
- Read data from a plain text file, which can be done using the code below (the file should be a plain text file from the same directory as your Notebook):

- o `f = open('file.txt', 'r')` # open a file for reading
`text = f.read()` # stores the text of the file in string `text`
`f.close()` # close the file

2. Complete any **three** of the following [25 points each]

- a. Generate a table showing the 20 most frequently used words and generate a word cloud. Stopwords should be removed, and all words should be converted to lowercase. You may also choose to *stem* or *lemmatize* your text, and may choose to restrict the word cloud to, e.g., the top 50 words if it is too cluttered.
- b. Repeat the above using *n*-grams. An *n*-gram corresponds to a set of *n* consecutive words. A *bigram* is an *n*-gram where $n = 2$ and a *trigram* is an *n*-gram where $n = 3$. For example, the sentence “how are you?” has bigrams of “how are” and “are you”. Using *textblob*, *n*-grams can be accessed using *blob.ngrams(n)*. Generate a table and/or wordcloud showing the most common bigrams and/or trigrams.
- c. Sentiment analysis (select at least one)
 - i. Output the five most positive sentences (or lines), and the five most negative sentences (or tweets).
 - ii. Compare the sentiment of two different text collections (e.g., two different Wikipedia pages, two different chapters, etc). Recommended: compare sentiments visually using pandas boxplot (<https://wellsr.com/python/python-create-pandas-boxplots-with-dataframes/>)
- d. Use named entity recognition to output the most commonly mentioned entities for various types, such as the most common people (PERSON), organizations (ORG) or geopolitical entities (GPE). Hint: use a *Counter* object to count the frequency of each element in a list (for an example see <https://www.hackerrank.com/challenges/collections-counter/problem>)
- e. Other analyses may be acceptable. I encourage you to be creative and contact me with ideas. I am also happy to help with data preparation. For example for a movie script, you may wish to compare the dialogue of two characters. In that case I can provide code for you to separate the dialogue for two characters.